

INTERNSHIP PROGRESS REPORT (Form 2-2)


(To Be Filed at the Completion of 1/2 of the Internship Period)

INTERN'S STATEMENT


Intern's Name: Najeeb Abdi


1. Describe what you did during this period.

Database schema (`src/db_models.py`) — one **Plan** owns its own **Course**/**AppConfig** rows so multiple curricula can be stored side by side without their data colliding; **User.active_plan_id** makes plan selection per-user rather than a single global mutable baseline.




Dashboard — the program roadmap animates term by term as the simulation runs, showing seat usage and status per course, colored by requirement type, with a running narrative of what happened each semester.





Live Simulation — a separate stepwise mode that advances one term at a time so an admin can review a term and edit capacity/policy knobs before advancing.



This period was spent building a discrete-term, agent-based simulation of Qatar University's Computer Science curriculum: modeling how students move through prerequisite chains, registration priority, and course capacity over up to 12 semesters, with the goal of identifying which bottlenecks cause delay or non-completion.

I chose an agent-based design over a simpler aggregate/cohort-flow model because the research question depends on each student's individual completed-course history — which prerequisites they personally satisfy, their own GPA and attempt count — not just cohort-level totals. Tracking per-student state was a requirement, not a convenience.

Week 1: Core Simulation Engine

I set up the project and built the core Python simulation engine from scratch:

- A discrete-term loop that advances student agents through registration, seat allocation, and pass/fail outcomes each semester.
- A multi-cohort model where a new cohort is admitted every year and all cohorts compete for one shared pool of course seats.
- GPA-driven dropout logic and capacity-bottleneck tracking for the CS core sequence.
- A `random.Random(seed + student_id)` stream per student, so a run is fully reproducible and one student's result never depends on the order other students are processed in — this matters later for attributing a change in results to a specific config change rather than random noise.
- Initial documentation (technical design doc and README) describing the model's assumptions.

Week 2: Full-Stack Dashboard

I moved from a script-only engine to a full product:

- Built a FastAPI backend on top of the simulation engine, a database layer (SQLAlchemy/SQLite) with JWT authentication, and a Next.js/React frontend dashboard. I chose FastAPI + SQLAlchemy specifically so the same simulation engine could be called identically from a CLI script and from an HTTP endpoint, with no duplicated logic between them.
- On the frontend: an animated curriculum graph that plays back the simulation term by term, the existing static analytics figures ported into React/SVG components, and a prerequisite-network diagram.
- Added a Scenario Builder for what-if experiments, a Plan Builder wizard for creating alternate curricula from scratch, a Settings page with full CRUD on courses/instructors, a capacity-planning module checking both seat and instructor staffing feasibility, a light/dark theme with a visual design pass, and a deployment configuration for hosting the app on Render.

Challenge: generalizing the engine from a hardcoded Fall/Spring cycle to also support optional Winter/Summer terms broke a formula I had written in Week 1. A student's simulation horizon was computed as `entry_term + max_terms`, which is only correct if every term counts toward a student's semester budget. Once optional terms existed, that formula truncated the simulation window and could cut a student off before they had actually used all of their real (mandatory-term-only) semesters.

I fixed it by replacing every call site with a dedicated `mandatory_horizon_end_term()` helper that walks the calendar counting only mandatory terms, and by giving each student a stateful `personal_semester` counter that only increments on a mandatory term, instead of recomputing it from `current_term - entry_term`. I caught the bug by adding a scenario with optional terms enabled and comparing enrollment before and after.

Week 3: Live Simulation + Scope Simplification

I added a "Live Simulation" mode that runs one semester at a time instead of the whole window at once, so an admin can review a term's results, adjust capacity/policy knobs, and manually advance to the next term —

closer to how a real academic-planning workflow works.

Challenge: my first instinct was to advance the live simulation by mutating the simulator's in-memory state one term at a time and saving a snapshot after each step. Partway through, I realized this made it impossible to guarantee that an edit made now (e.g., adding a course section) would never retroactively change a term that had already been reviewed and saved earlier — which breaks the whole point of a term-by-term review tool.

I re-architected it around a pure replay model instead: every "advance" call replays the engine from term 0 with the full, current list of edits (each tagged with the term it takes effect from), so a term already saved is guaranteed to reproduce byte-identically no matter how many later edits are added. I verified this with an automated test that adds an edit after the fact and asserts the earlier snapshot's frame is unchanged, rather than just checking it by hand.

After that, I did a significant simplification pass. After reviewing what was actually load-bearing for the research question, I removed authentication, the Scenario Builder, the Capacity Planning page, and the instructor-staffing model to cut the product down to its core value, and added an inline "what-if" panel plus capacity-section recommendations directly on the main dashboard.

Week 4: UI Refactor

I refactored the dashboard's navigation and component structure based on usability issues I noticed while using the tool myself:

- Extracted a dedicated pre-start screen.
- Moved the what-if panel into the Bottlenecks page and promoted Bottlenecks to primary navigation.
- Added live running totals to the in-progress simulation view.

2. Describe what you learned during this period.

I learned how to design and build a complete simulation product end-to-end: from the modeling logic (agent-based simulation, prerequisite/eligibility rule evaluation, seat-allocation policy) through a REST API layer to an animated frontend dashboard.

- **Backend:** how to structure a FastAPI service with a proper database layer, authentication, and a clean separation between the simulation engine and file I/O, so the same engine could be called from both a CLI script and an API.
- **Frontend:** how to build animated, data-driven visualizations in React/SVG and how to lay out a dashboard so a non-technical user (an academic-planning admin) could actually make sense of the output.
- **Concretely new to me:** I did not previously know how to reason about a "replay to guarantee determinism" pattern (Week 3's Live Simulation challenge above); I now understand it as a general technique for any system that needs to let past state be revisited safely while still accepting new input going forward.

Connection to my coursework. This project repeatedly pulled on specific things I studied at QU rather than generic "programming practice":

- **Data Structures and Algorithms** — the prerequisite/eligibility rules are modeled and validated as a directed graph with an explicit cycle check ([networkx](#)-based). I had to actually reason about what a

cycle in a prerequisite graph would mean for the simulation (a course that can never become eligible) rather than just calling a library function.

- **Database Systems** — the schema, where one **Plan** owns its own **Course/Config** rows so several curricula can be stored side by side without their data colliding, is a direct application of relational schema design and normalization.
- **Software Engineering** — the engine/API/CLI separation (**service.py** as a pure, file-I/O-free boundary called by both **run.py** and the FastAPI layer) is the layered-architecture and separation-of-concerns principle, applied to a real, growing codebase instead of a textbook example.
- **Probability and Statistics** — the reproducible, seeded-RNG-per-student design, plus the Monte Carlo confidence intervals reported alongside headline metrics, draw directly on the coursework covering sampling variance and confidence intervals.

I also learned a project-management lesson I didn't expect going in: after building out a feature-rich version of the tool (auth, multi-plan support, capacity planning, instructor modeling), I realized a lot of that complexity wasn't serving the core research question, and I learned to make the call to strip it back down rather than keep adding features. That simplification pass was one of the most useful things I did this period; it taught me to keep re-asking what the tool is actually for, instead of just accumulating functionality.

Working mostly independently on a long-running codebase also pushed me to get better at reading my own prior design decisions back (via documentation I'd written myself) and staying consistent with them as the project grew, which is a different skill than writing code for a one-off assignment.

Going forward, the concrete change this makes to how I work: before adding a feature, I will now explicitly write down which research/product question it serves, the same way I did retroactively during the Week 3 simplification pass, instead of only doing that check after the complexity has already piled up. I will also default to a replay-from-source-of-truth design (rather than incremental in-place mutation) any time a system needs to support "go back and change something earlier," since debugging the Live Simulation determinism issue cost me more time than designing it in from the start would have.

Intern's Name: Najeeb Abdi **Date:** 04/07/2026

EMPLOYER'S STATEMENT

Please comment on the intern's work during this period.

Industrial Supervisor's Signature: _____ **Date:** _____